

# Few-Shot Learning for Text Classification

Nowcasting and Forecasting with Text as Data

---

Renato Vassallo

Institute of Economic Analysis (IAE-CSIC) - Barcelona School of Economics

May, 2025

1. Introduction
2. Embeddings and LLMs
3. Zero-Shot Learning
4. Few-Shot Learning
5. Applications

# Introduction

---

# About Me

## Academic Background

- B.A. in Economics - UDEP (2016)
- M.A. in Economics - PUCP (2021)
- M.Sc. in Data Science - BSE (2023)

## Contact Information

- ✉ : [renato.vassallo@bse.eu](mailto:renato.vassallo@bse.eu)
- in : [linkedin.com/in/renatovassallo](https://www.linkedin.com/in/renatovassallo)
- 🐙 : [github.com/RenatoVassallo](https://github.com/RenatoVassallo)

**Goal:** share what I know and make it useful for you.



1. **Block 1:** Concepts and code walkthrough.

- Github repo: [here](#).

2. **Block 2:** Hands-on activity. [▶ Go to Activity](#)

- Select a topic and apply any method discussed in Block 1.
- Work in groups (max 4 members).
- Duration: 30-45 minutes.
- Brief presentation: 5 minutes per group.

# Important notice


- The class materials, code, and methods are designed for teaching and illustrative purposes only.
- We will work with large Transformer models, so access to GPUs is highly recommended for optimal performance.
- You are encouraged to improve, extend, and optimize the routines — consider implementing parallelization where applicable.

## Session 1 - Objectives

- Transforming text into numerical representations.
- Fine-tuning large language models (LLMs) for specific tasks.
- Leveraging state-of-the-art methods for text-based indicators.
- Real-world applications: Industry and academia.

*How can we extract meaningful patterns from text data?*

# Semantic coverage






# Embeddings and LLMs

---

# What are embeddings?

*When dealing with text data, how do we convert words and sentences into something that a model can understand?*

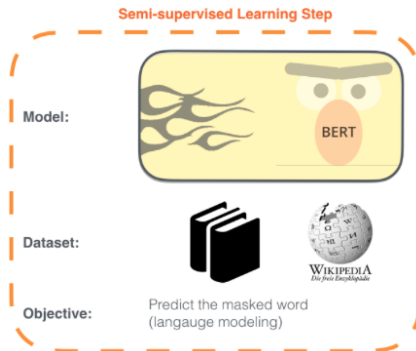
- Embeddings are dense vector representations of text.
- They capture semantic information in a continuous space.
- The closer two vectors are, the more similar the sentences are in meaning.



Source: Mayoral et al. (2025).

# How embeddings are generated?

- BERT (Devlin et al., 2019) is pre-trained on large corpora.
- With pre-trained BERT we can directly obtain embeddings for any text.



Source: Alammari (2019).

# BERT embedding similarity

Let's see how we can use BERT to quickly assess the similarity between two sentences:

```
from transformers import AutoTokenizer, AutoModel
import torch

tok = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModel.from_pretrained("bert-base-uncased")

emb1 = model(**tok("Cats are cute", return_tensors="pt")).last_hidden_state[0, 0]
emb2 = model(**tok("Dogs are loyal", return_tensors="pt")).last_hidden_state[0, 0]

print(torch.cosine_similarity(emb1, emb2, dim=0).item())
```

- But many NLP tasks like sentiment analysis, NER, Q/A, or text classification require task-specific models.
- BERT embeddings are generic; we can fine-tune a model to adapt to specific tasks.
- Let's look at how to train our own simple LM for text classification!

► [Go to Training simple LM](#)

# Scaling up: efficient alternatives to LLMs

- Ideally, we could leverage state-of-the-art models like OpenAI GPT-4, Anthropic Claude, or Google PaLM 2, but unlimited access is costly and resource-intensive.
- On the other hand, training large models from scratch requires substantial computational power, time, and expertise.
- Fortunately, **pre-trained models** provide effective starting points, tailored for various NLP tasks:
  - BERT Variants - fine-tuned BERT for specific tasks like sentiment analysis or NER.
  - Hugging Face Models Hub - extensive library of models for NLI, NER, Q/A, and more.
  - Cohere Models - focused on semantic understanding and text generation.

# Sentence transformers

- Sentence Transformers: efficient embeddings for semantic search, clustering, etc.
- Lightweight models can provide a balance between performance and efficiency.

Inference Time Comparison for 100 Texts

Model	Time (seconds)
BERT	0.5728
SBERT	0.4985
all-MiniLM-L6-v2	0.0900

- **BERT**: Larger and slower due to more layers and parameters. Layers: 12; Embedding size: 768.
- **SBERT**: Optimized for sentence-level embeddings but still relatively heavy.
- **all-MiniLM-L6-v2**: Much smaller and faster, designed for speed without substantial accuracy loss. Layers: 6; Embedding size: 384.

# Zero-Shot vs. Few-Shot Learning

1. **ZSL:** make predictions without having seen specific labels during training.
  - **Example:** predict the probability of a news headline referring to a coup d'état.
2. **FSL:** adds a **supervised layer** with a small set of labeled examples in a specific domain.
  - **Example:** predict the probability of a news headline referring to a coup d'état, given a few labeled examples of similar headlines.

	Zero-Shot	Few-Shot
<b>Training data</b>	None for specific labels	Few labeled examples
<b>Flexibility</b>	Broad/generalized	More specific
<b>Accuracy</b>	Reasonable	Higher (for seen classes)
<b>Efficiency</b>	Faster	Slower (additional training)



# Zero-Shot Learning

---

# What is Zero-Shot Learning (ZSL)?

- **Definition:** Zero-shot learning enables models to classify data without having seen specific labels during training.
- **How it works:**
  - Leverages semantic information (e.g., embeddings) to link known and unknown classes.
  - Uses pre-trained models like BERT, GPT, or SBERT to generate rich contextual embeddings.
- **Example:** If trained on animals (cats, dogs) but asked to classify birds, the model uses textual similarity to infer the correct class.

## References:

- *A comprehensive review on zero-shot-learning techniques* (Lazaros et al., 2024)
- Hugging Face Zero-Shot Pipeline: [link](#)

## 1. Pre-trained model as encoder

- BERT, SBERT, GPT, etc., are used to generate embeddings for input text.

## 2. Label embedding generation

- Transform labels into descriptive sentences or keywords.
- Example: instead of “Positive”, use “This text expresses positive sentiment”.

## 3. Similarity scoring

- Compute cosine similarity between input text embedding and each label embedding.
- Assign label based on the highest similarity score.

**Goal:** map documents to labels using similarity in embedding space.

- Given a query  $q$  (e.g., a news article), predict a label  $\ell$  from a set  $\mathcal{L}$ .
- Apply an encoder to both the query and each label to obtain embeddings.
- Assign the label that maximizes similarity with the query embedding in latent space.

The predicted label  $\hat{\ell}$  is defined as:

$$\hat{\ell}_{ZSL} = \arg \max_{\ell \in \mathcal{L}} \cos(\mathbf{E}(q), \mathbf{E}(\ell))$$

# Natural Language Inference (NLI)

- NLI considers two sentences: a *premise* and a *hypothesis*. The task is to determine whether the hypothesis is true (**entailment**) or false (**contradiction**) given the premise.

Premise	Label	Hypothesis
The cat is sleeping on the couch.	Contradiction	The cat is playing outside.
The company reported a rise in profits.	Neutral	The company launched a new product.
A group of people is protesting in the street.	Entailment	There is a protest happening.

Examples of NLI: Contradiction, Neutral, and Entailment


► [Go to NLI Architecture](#)

- We will use pre-trained NLI models like `bart-large-mnli` and `roberta-large-mnli`.
  - These models are trained on large datasets like MNLI (Multi-Genre Natural Language Inference), containing over 430k pairs of sentences labeled as **entailment**, **neutral**, or **contradiction**.
- Methodology:
  - Take the input text as the **premise**.
  - Convert each candidate label into a natural language **hypothesis**.
  - The NLI model predicts whether the premise **entails**, is **neutral**, or **contradicts** each hypothesis.
- Interpretation:
  - If the model predicts **entailment**, we consider the label as likely true.
  - This method allows for classification without explicit training data for each label — a key strength of zero-shot learning.

# Simple zero-shot classifier using NLI

```
text = "Opposition party gains ground ahead of national election."  
labels = ["economics", "sports", "politics"]
```

```
ZeroShotNLI(tokenizer, model, text, labels, plot=True)
```



Output:

```
{'sequence': 'Opposition party gains ground ahead of national election.',  
 'labels': ['economics', 'sports', 'politics'],  
 'scores': [0.028, 0.017, 0.955]}
```

# Few-Shot Learning


---



# Why few-shot learning?

- There is a need to develop models that can generalise with limited data.
  - Medical diagnosis with few patient records.
  - Image classification with few labels (detecting deforestation, face recognition).
  - Rare event detection in news or official statements.
- Few-shot = Zero-shot + Supervised Layer.
  - Supervised layer: Fine-tunes with a small number of labeled examples.
- Be careful with **overfitting**!
  - Too few examples can lead to memorization rather than generalization.
  - Regularization techniques are crucial.

# Few-shot flowchart



Source: Agrawal (2024), via LinkedIn.

# A simple approach

**Goal:** learn a mapping between the documents and their labels.

- Weights, ( $\mathbf{W}$ ), are learned through minimizing the loss function, as expressed below:

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} (\|\mathbf{X}^\top \mathbf{W} - \mathbf{Y}\|^2 + \lambda \|\mathbf{W} - \mathbb{I}\|^2)$$


- 1st term:* tells  $\mathbf{W}$  how to map  $\mathbf{X}$  to  $\mathbf{Y}$ . *2nd term:* the elements of the weight matrix are pushed towards the identity matrix.
- If few examples:  $\mathbf{W}$  will likely be quite close to the identity matrix (FSL  $\sim$  ZSL).
- If many examples:  $\mathbf{W}$  will be pushed further away from the identity matrix.

The predicted label ( $\hat{\ell}$ ) is obtained by maximizing the *similarity* between the query embedding and each label embedding, as follows:

$$\hat{\ell}_{FSL} = \arg \max_{\ell \in \mathcal{L}} \cos(\mathbf{E}(q)\mathbf{W}^*, \mathbf{E}(\ell)\mathbf{W}^*)$$

# How many labels are enough?

- **AG News:** subset of 120k for training and 7.6k for testing. 4 categories.
- **Reddit:** subset of 540k for training and 60k for testing. 16 categories.



Source: Cloudera (2020).

# Applications

---

# Applications

---

## Training a simple LM

- We use a simple dataset of sentences with categories: Learning, Pets, Coding.
  - ("I love machine learning", [1, 0, 0])
  - ("Cats are cute", [0, 1, 0])
  - ("Python is great for programming", [0, 0, 1])
- Model Architecture:
  - Pre-trained BERT model as the encoder.
  - A linear layer to project the CLS token to a 3D space.
  - MSE Loss to align embeddings with target vectors.

*Original  
sentence*

**“Cats are cute”**

- **Input Text:** We start with a simple sentence. This is just a string of words, but the model can't understand text directly.
- We need to convert it to numbers.



# Training Process

- **Tokenization:** breaks down the sentence into individual words or subwords and assigns each a unique number (ID).
- These numbers are packed into a fixed-size array. This ensures that every input has the same length.

*Original sentence*

"Cats are cute"



*Tokens*

[CLS]

Cats


are

cute

[SEP]

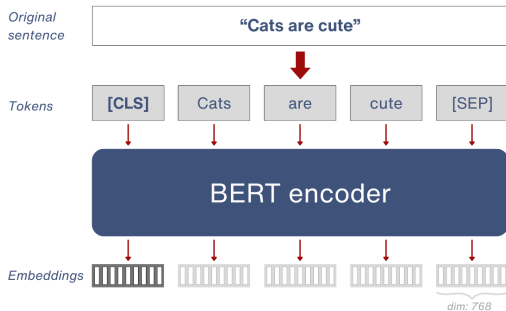
# Training Process

- **BERT:** Think of BERT as a huge library of knowledge about language.
- It has been pre-trained on massive amounts of text to understand grammar, meaning, and context.



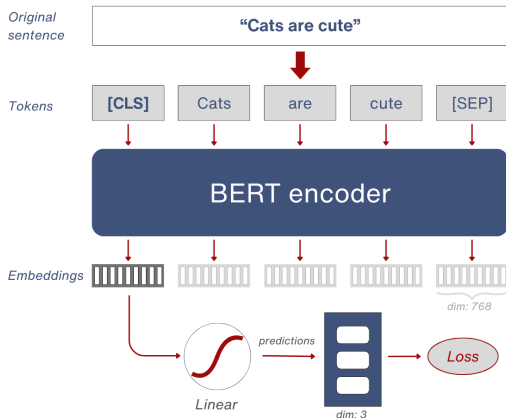
# Training Process

- Takes in the tokenized numbers and converts them into meaningful vectors that capture the context of the words.
- **Example:** The word *Python* could mean a programming language or a snake.
- BERT can understand the context and represent it as a vector that reflects the intended meaning.

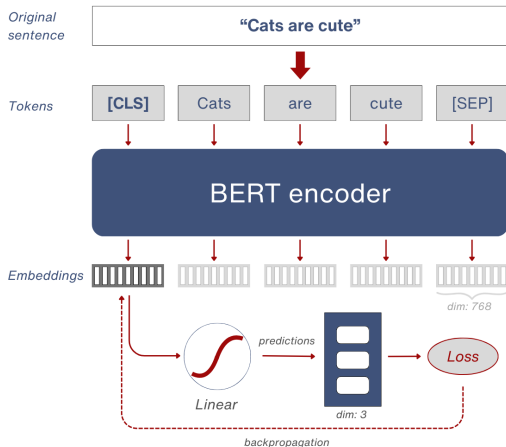


# Training Process

- BERT outputs a vector of 768 dimensions for **each token**.
- The **linear layer** acts as a projector, reducing those 768 dimensions down to just **3 dimensions**.
- For *Cats are cute*, the target vector is  $[0, 1, 0]$ .
- The loss function defined in the model is Mean Squared Error (MSE).



- The computed loss is then **backpropagated** through the model.
- The optimizer updates the parameters of the linear layer and potentially some of the BERT parameters.




# Applications

---

## Natural Language Inference

- NLI datasets are typically modeled via *sequence-pair classification*.
- The input premise and hypothesis are encoded (e.g. using BERT). The obtained vectors  $u$  and  $v$  are then concatenated along with their element-wise product and absolute difference.
- This representation captures information from both inputs. This vector is then passed to a 3-class classifier consisting of multiple fully-connected layers.



Source: Sadeghi et al. (2022).

# Applications

---

## Hands-on activity



# Mini Project 1: political violence analysis

**Data:** ACLED, daily data.

- **Use case:** analyze the increase in riots and protests in Sudan. Characterize these events as peaceful or violent.
  - Visualizations: time series plots, heatmaps, geographical maps.
- **Predictive task:** divide data into train (until Dec 2021) and test (2022 onwards). Train a classifier to predict event types (e.g., riots/protests vs battles, explosions).
  - Products: ROC-AUC, PR-AUC, Confusion Matrices.
- **Open exploration:** identify other countries or event types. Use embeddings to extract patterns, but keep in mind computational limits.
- **Support:** if you need specific data for a country, ask and I can retrieve it.

## Mini Project 2: Banking77 dataset analysis

**Data:** Banking77, intent classification dataset.

- **Use case:** analyze customer support requests to identify common intents such as fraud reporting, balance inquiry, or card blocking.
  - Visualizations: distribution of intents, time-based trends, sentiment analysis.
- **Predictive task:** train a few-shot learner to classify intents based on support requests.
  - Products: confusion matrices, precision-recall curves.
- **Open exploration:** use embeddings to cluster similar intents or identify misclassifications.
- **Support:** you can subset data to focus on specific intents to reduce training time.

**Goal:** extend the application of Lab 2 and apply more advanced techniques for economic analysis.

- Use few-shot learning and compare its performance against zero-shot predictions.
- Replace the AR(1) model with your favorite ML regressor (RF, CatBoost, etc).
- Apply a rolling forecast with TimeSeriesSplit to prevent data leakage.
- Visualize predictions and compute error metrics.

# Thank You

## Questions?